

# CanvasGPT

Paris Phan, Ethan Zhang

Nov 11, 2025

# Background / Motivation

- Canvas offers educators a high degree of freedom in structuring their courses
  - Syllabus
  - Announcements
  - Grades
  - Gradescope
- Every teacher has a canvas layout that differs slightly
  - Some use announcements for assignment details instead of 'assignment' tab
  - Some separate assignment submissions from details, incorporating Gradescope
  - Some teachers use their own course website (e.g. github page or Webassign) for all information, and canvas becomes just a proxy
- Students have trouble keeping track of course structure and materials across 5+ classes in addition to possible external sites
- It can be difficult to ensure you're on top of everything if you're unable to find specific information

UVA

CS\_4444-001 > Modules

2025 Fall

Home

**Modules**

Assignments

Grades

UVA Library Portal

Piazza

Gradescope

Online Meetings

Groups

Calendar

Inbox

History

Studio

Commons

My Media

Help

Lecture Slides

- Week-1-updated.pdf
- Week-2-1.pdf
- Week-2-2.pdf
- Week-3-updated.pdf
- Week-4-quiz.pdf
- Week-4.pdf
- Week-5-1.pdf
- Week-5-2.pdf
- Week-6.pdf
- Week-7-revision.pdf
- Week-8.pdf
- Week-9-1.pdf
- Week-9-2.pdf

Collapse all

canvas.its.virginia.edu

UVA

PSYC\_2150-001

2025 Fall

Home Announcements Syllabus Modules Assignments Quizzes Grades Files 1 SensusAccess Course Evaluations Online Meetings UVA Bookstore Inclusive Access Course Email People UVA Library Portal Collaborations Posted Feedback Anonymous Feedback Panopto Video

Switch to Old Files Page All my files

## Files

Search files

Enter at least 2 characters to search

25F Introduction to Cognition

0 selected

	Name	Created	Last Modified	Modified by	Size	Status	Actions
<input type="checkbox"/>	Alternate Files	14 Oct 2025	14 Oct 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Participant Pool	25 Aug 2025	25 Aug 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 1	1 Nov 2022	26 Aug 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 2	2 Feb 2024	4 Sep 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 3 (Research Methods)	2 Feb 2024	12 Sep 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 4 (Visual Perception)	9 Feb 2024	15 Sep 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 5 (Exam 1)	10 Feb 2025	22 Sep 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 6 (Attention)	21 Feb 2025	3 Oct 2025	--	--		<input type="button" value="..."/>
<input type="checkbox"/>	Week 7 (Systems of Thought and Biases)	1 Nov 2022	6 Oct 2025	--	--		<input type="button" value="..."/>

# Related Work



Featured 3.9 ★ (34 ratings) Share

Extension

Workflow & Planning

70,000 users

- Offers slight convenience features
- No AI integration
- No Smart Assistance
- No Information compilation

## Search Bar.



## Related Work

◆ Gemini

⋮ ⌛ ✕

See what Gemini can do >

What can Gemini do  
in Google Drive

Summarize a topic  
based on files in my Drive

Summarize a folder  
in my Drive

Gems >

M Mathematical Statistics

Sales pitch ideator  
Create compelling sales pitches that resonate with your audience and...

Learning coach  
Here to help you learn and practice new concepts. Tell me what you'd...

Your Gems will appear across Workspace

Massive amount of precedence in work platform AI indexing

- Many workspace platforms integrate AI indexing and search
  - Google Drive
  - Notion Workspace AI
  - Slack AI
  - Microsoft Copilot

Even local document-based assistants

- Warp (AI Terminal)
  - Claude Code
  - Cursor
  - etc.

# Claim / Target Task

To make life easier for college students with an interface for Canvas that unifies, normalizes, and semantically indexes course content and assignments across inconsistent instructor setups, enabling intelligent retrieval and proactive alerts for students.

# Proposed Solution

An app / platform that connects to your Canvas student account, that utilises AI/ML to

- Read and understand all your courses and their individual unique structures
- Ingest and unify data to provide a single structure to view all your course information
- Automatically detects new updates to courses, including new assignments/grades/etc.
- Notifies you of new updates, upcoming deadlines, etc.
- (backlog) Integrates with your workspaces (Google drive, notion, etc.)
  - Automatically updates and fills your workspace with new material

# Implementation

## Electron Desktop App

- Utilising local ML models + API LLMs

## 3 Major Components

- LangGraph Discovery Agent
- Unstructured Data Discriminator
- MCP LLM Interface

# LangGraph Discovery Agent

Agentic workflow to automatically discover  
the routes of data sources for future data ingestion

- Starts at the canvas home page
- Determines valuable native canvas endpoints
- Explores linked external sites for missing sources

Needed Data Sources

Syllabus

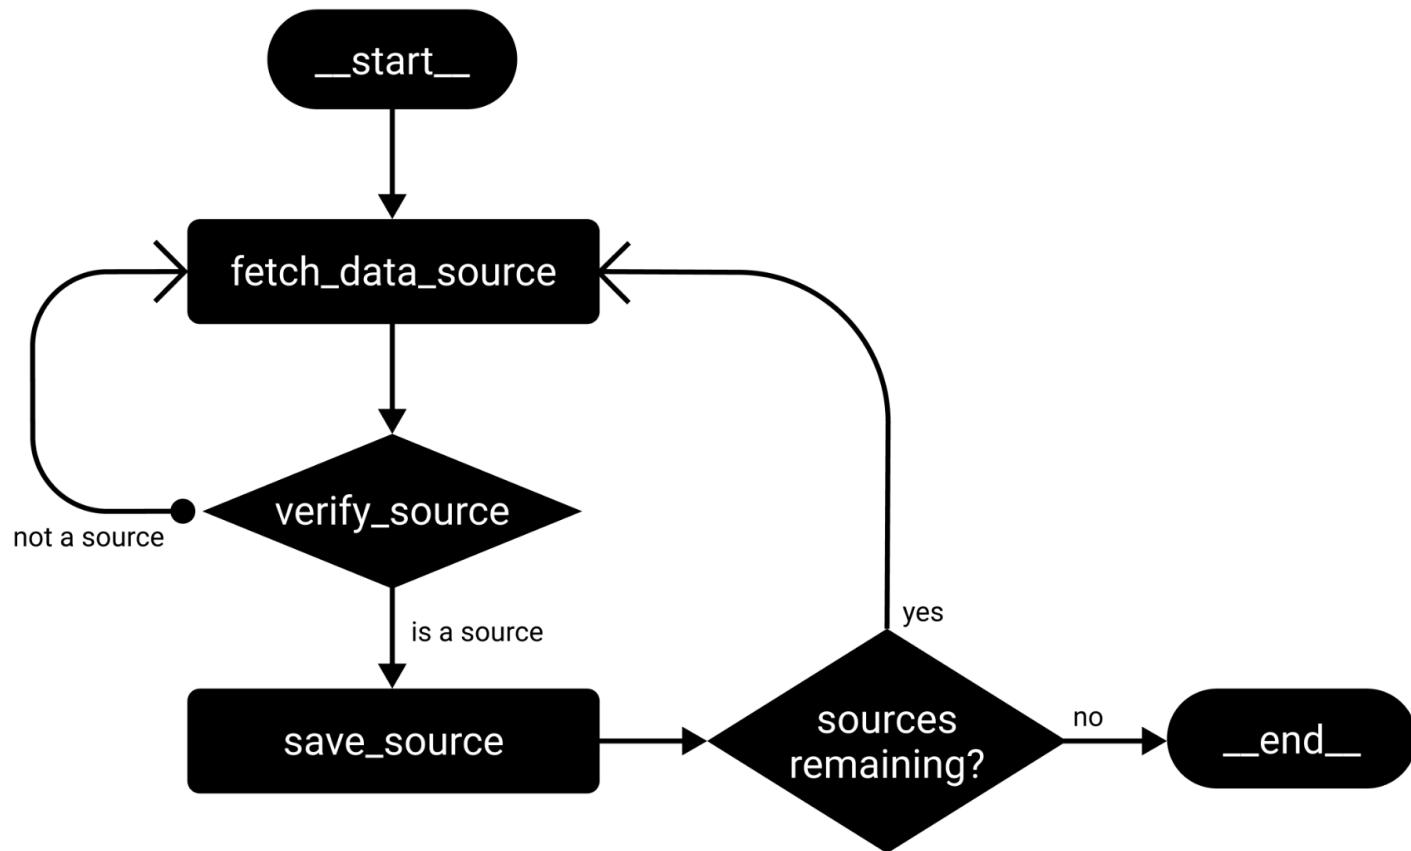
Assignments

Course Learning Material

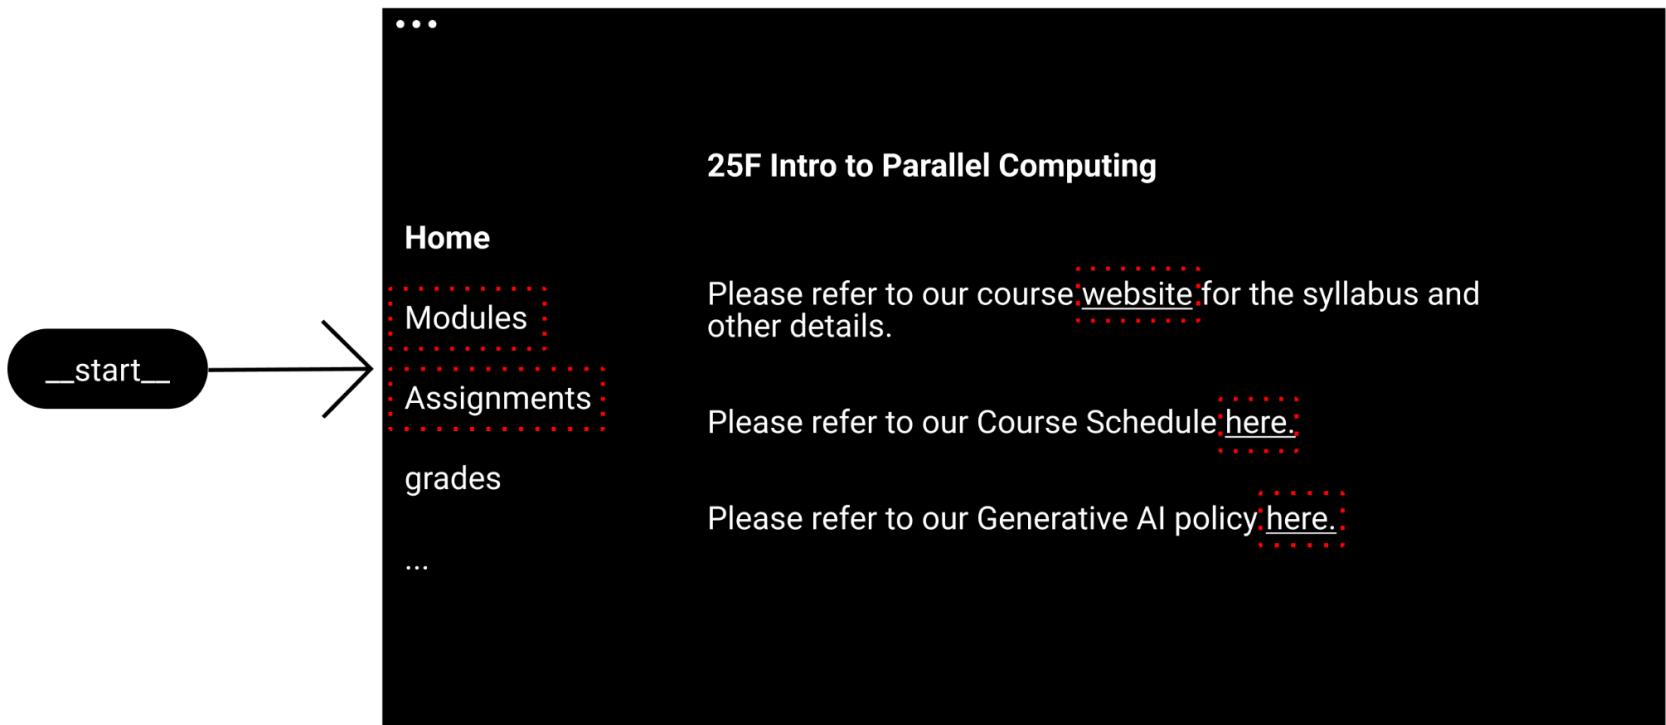
Schedule

```
1 // --- Graph Definition ---
2
3 const graph = new StateGraph(DiscoveryState)
4   .addNode("fetch_data", fetchCourseData)
5   .addNode("analyze_structure", analyzeStructure)
6   .addNode("verify_source", verifySource)
7   .setEntryPoint("fetch_data")
8   .addEdge("fetch_data", "analyze_structure")
9   .addEdge("analyze_structure", "verify_source")
10  .addEdge("verify_source", END);
```

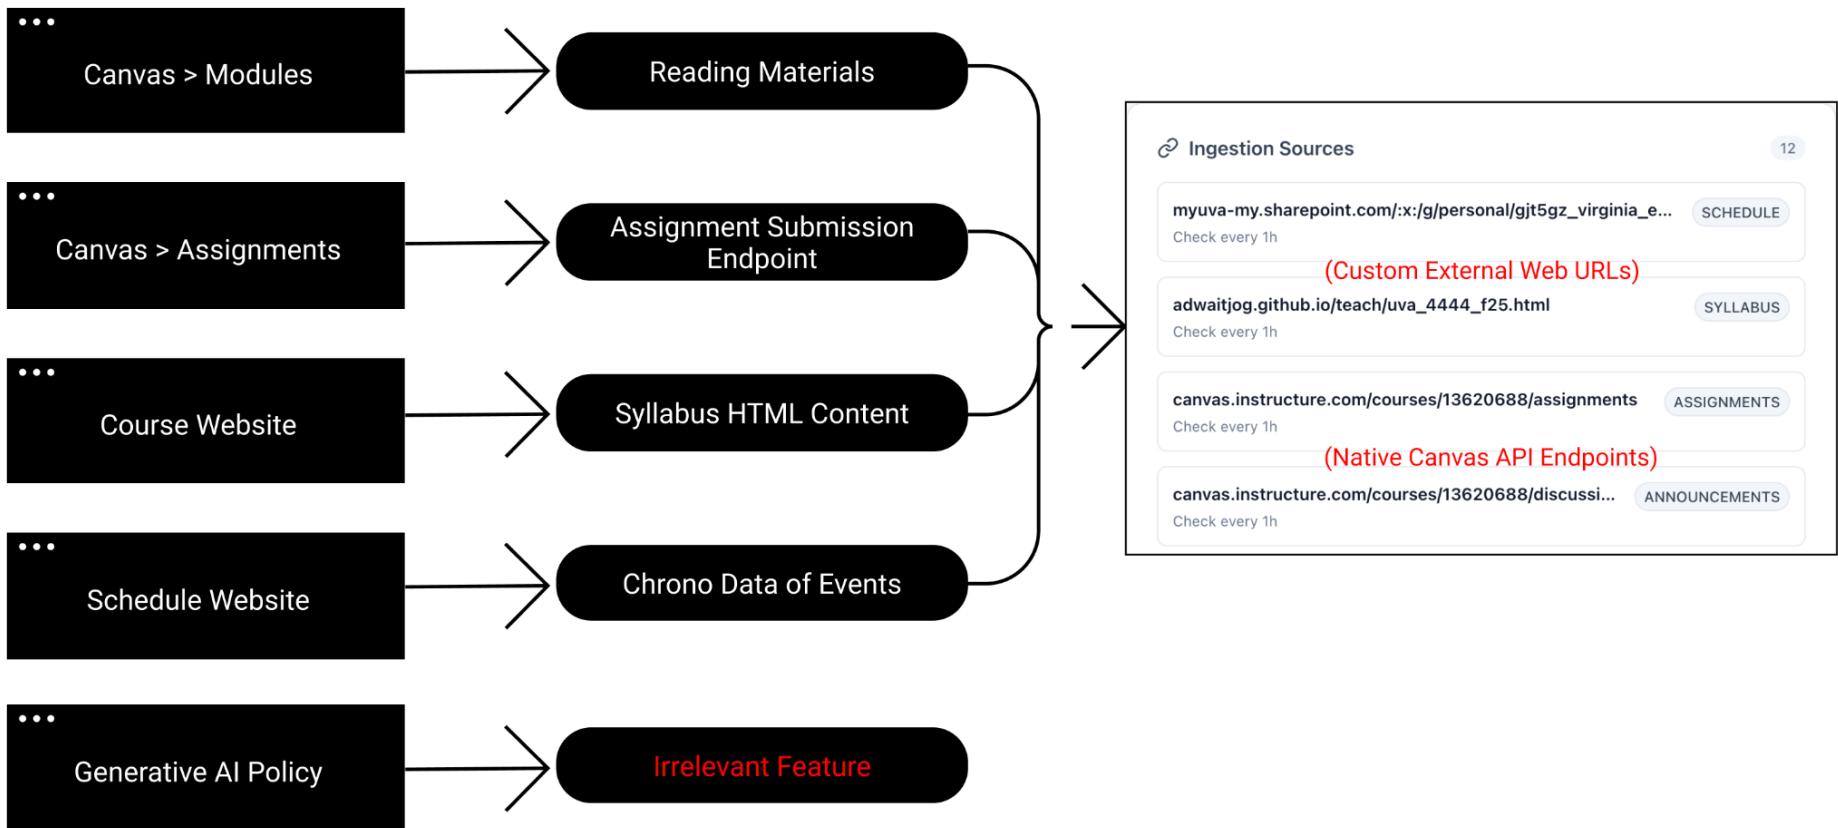
# LangGraph Discovery Agent



# LangGraph Discovery Agent



# LangGraph Discovery Agent

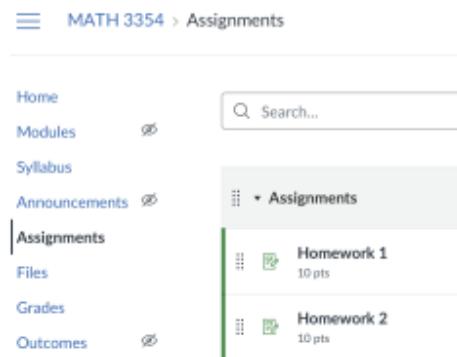


# Data From Ingestion Sources

We have all possible data for a course

BUT, they are all unorganised and unstructured

For example:

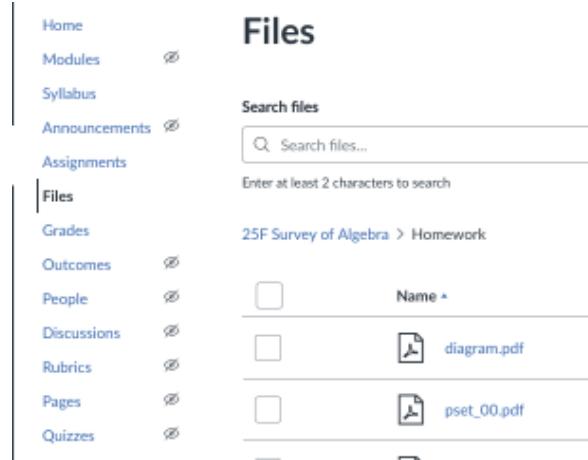


The screenshot shows the Canvas interface for a course titled "MATH 3354". The left sidebar menu includes "Home", "Modules", "Syllabus", "Announcements", "Assignments" (which is currently selected), "Files", "Grades", and "Outcomes". The main content area is titled "Assignments" and lists two assignments: "Homework 1" (10 pts) and "Homework 2" (10 pts). A search bar is at the top right.

Instead, this information  
is found in a pdf, in  
Canvas > Files!

Canvas > Assignments

Gives us the submission  
endpoint for our assignments,  
but **no due date, no task info**



The screenshot shows the Canvas interface for a course titled "25F Survey of Algebra". The left sidebar menu includes "Home", "Modules", "Syllabus", "Announcements", "Assignments", "Files" (which is currently selected), "Grades", "Outcomes", "People", "Discussions", "Rubrics", "Pages", and "Quizzes". The main content area is titled "Files" and shows a list of files: "diagram.pdf" and "pset\_00.pdf". A search bar and a placeholder message "Enter at least 2 characters to search" are also present.

# Unstructured Data Discriminator

Each piece of data is stored as a ‘universal item’ in our database

- Highly unstructured, no item rules
- [type: Assignment, name: Homework 1, metadata: NULL]
- [type: File, name: pset-1.pdf, metadata: {due\_date, desc, ...}]

We need to find a way to connect these two pieces of data and form one true **realised entity** to show the user.

- Confident in its existence
- Confident in its metadata (correct due date)

# Unstructured Data Discriminator

## Phase 1: Feature Extraction & Vector Embedding (?)

- extract ‘assignment number’ from every string using regex

## Phase 2: Clustering Algorithm

- group all related items together (e.g. pset01.pdf, homework 1)

## Phase 3: Verify confidence levels

- Prompt user to ‘triage’ items with confidence  $< 0.9$

# MCP LLM Interface

## 3 Major Components

- MCP LLM Interface
- Host/Expose the Electron App as a MCP Server
- Claude Code is able to access this MCP server, calling the exposed tools
  - Gives Claude the context of all the canvas ingestion
    - Can help with Homework
    - Plan out homework
    - etc.
  - Will extend to other tools (Cursor, Notion, Zapier, IFTT, etc.)

# Challenges

## Unstructured Data Discrimination Model

- Rule Based Fuzzy Matching Algorithm
  - Lacks semantic understanding for unstructured courses
- Discriminative Entity Recognition (Random Forest, XGBoost))
  - Requires training dataset
- Feature Extraction and Clustering
  - Timeline Feasible

# Future Improvements

Unstructured Data Discrimination Model

Hybrid Graph-Based Resolution Approach

1. Feature Extraction & Blocking
  - Rule matching & Normalisation
2. Candidate Generation
  - SBERT Semantic Embeddings to compute cosine similarities
3. Transitive Graph Resolution
  - Union-Find to find connected components

# Application Package

- Electron App (Electron-Vite)
- LangGraph
- openAI
- MCP
- Claude Desktop
- LanceDB
- TailwindCSS

# Any Questions?