

# Image Colorization

Sean Bremer & Aidan Szilagyi

December 15, 2025

# Motivation

## Problem

- Many archival images, historical photos, and creative projects are presented in grayscale
- Limits impact and accessibility

## Solution

- Colorization can restore and enhance images, making them more engaging and informative

# Target Task



- Infer realistic color channels from grayscale input
- Generative task

# Related Work

original




grayscale



Anwar et al. (2024) attempts to survey and benchmark various colorization methods

Larsson et al. (2016) Identified Five Common Failure Cases:



Too Desaturated

Inconsistent Chroma

Inconsistent Hue

Edge Pollution

Color Bleeding

# Proposed Solution

Representation	Score Function	Optimization	Model	Metrics
LAB Color Code Input: L channel  Output: A and B channels	L1 Loss Perceptual Loss GAN Loss	Stochastic Gradient Descent with Adam Optimization	Convolutional Encoder-Decoder. U-Net Style (Bottleneck, Skip Connections)	Structural Similarity Index Measure

## Representation (LAB Color Codes)

- Three axes color codes: greyscale, green-red, and blue-yellow
- LAB better models how the human eye perceives differences in color

## Model (U-Net Encoder Decoder)

- Encoder extracts low-level features, decoder reconstructs high-level features

## Optimization (Generative Adversarial Network)

- Generator takes grayscale input, generating colorized output
- Discriminator evaluates if the colorization is realistic, helping generator improve output

## Metrics (Structural Similarity Index Measure)

- Compares structural similarity between predicted and ground truth colorized images
- Accounts for luminance, contrast, and structure

# Implementation

ResNet-34 Encoder -> Our Decoder

Decoder:

Conv Channels: 512 -> 256 -> 128 -> 64

Each Conv has kernel size of  
channels/2 x channels/2

```
class ColorizationUNet(nn.Module):
    def __init__(self, pretrained=True):
        super().__init__()
        self.encoder = ResNet34Encoder(pretrained=pretrained)

        self.up4 = nn.ConvTranspose2d(512, 256, 2, stride=2)
        self.conv4 = nn.Sequential(
            nn.Conv2d(256 + 256, 256, 3, padding=1),
            nn.ReLU(inplace=True)
        )

        self.up3 = nn.ConvTranspose2d(256, 128, 2, stride=2)
        self.conv3 = nn.Sequential(
            nn.Conv2d(128 + 128, 128, 3, padding=1),
            nn.ReLU(inplace=True)
        )
```

Discriminator:

Conv Channels: 5 -> 64 -> 128 -> 256 -> 512 -> ReLU

Each Channel has 4x4 kernel with stride=2, halving the width and height each layer

```
class PatchGANDiscriminator(nn.Module):
    def __init__(self, in_channels = 5):
        super().__init__()

        self.conv1 = nn.Conv2d(in_channels, 64, kernel_size = 4, stride = 2, padding = 1)
        self.conv2 = nn.Conv2d(64, 128, kernel_size = 4, stride = 2, padding = 1)
        self.conv3 = nn.Conv2d(128, 256, kernel_size = 4, stride = 2, padding = 1)
        self.conv4 = nn.Conv2d(256, 512, kernel_size = 4, stride = 2, padding = 1)
        self.conv5 = nn.Conv2d(512, 1, kernel_size = 4, stride = 1, padding = 1)

        self.leaky_relu = nn.LeakyReLU(0.2, inplace = True)
```

# Data Summary


- Due to the large size of our training dataset, we trained our model on a randomized 10% of the dataset (roughly 370 images)
- For each epoch (10), we had the terminal print results every 100 steps (out of 370), depicting discriminator & generator loss (D & G loss)
- Sample run results:

```
Epoch [1/10], Step [1/370], D Loss: 1.3863, G Loss: 33.4629
Epoch [1/10], Step [101/370], D Loss: 1.2811, G Loss: 36.1519
Epoch [1/10], Step [201/370], D Loss: 1.3369, G Loss: 36.0646
Epoch [1/10], Step [301/370], D Loss: 1.1108, G Loss: 33.8154
Epoch [1/10], Step [370/370], D Loss: 1.2225, G Loss: 29.8015
Validation Results - L1 Loss: 0.0676, SSIM: 0.4171
Epoch [2/10], Step [1/370], D Loss: 1.2222, G Loss: 30.8697
Epoch [2/10], Step [101/370], D Loss: 0.6844, G Loss: 50.3742
Epoch [2/10], Step [201/370], D Loss: 0.8445, G Loss: 38.0917
Epoch [2/10], Step [301/370], D Loss: 0.8631, G Loss: 37.0352
Epoch [2/10], Step [370/370], D Loss: 1.0712, G Loss: 41.3892
Validation Results - L1 Loss: 0.0747, SSIM: 0.3219
Epoch [3/10], Step [1/370], D Loss: 0.9538, G Loss: 38.6591
Epoch [3/10], Step [101/370], D Loss: 0.7869, G Loss: 37.3652
Epoch [3/10], Step [201/370], D Loss: 0.6831, G Loss: 41.3719
Epoch [3/10], Step [301/370], D Loss: 0.5297, G Loss: 48.4916
Epoch [3/10], Step [370/370], D Loss: 0.8284, G Loss: 41.0376
Validation Results - L1 Loss: 0.0761, SSIM: 0.3259
Epoch [4/10], Step [1/370], D Loss: 0.7538, G Loss: 37.5349
Epoch [4/10], Step [101/370], D Loss: 0.6509, G Loss: 43.4072
Epoch [4/10], Step [201/370], D Loss: 0.7111, G Loss: 39.6706
Epoch [4/10], Step [301/370], D Loss: 1.1508, G Loss: 41.6996
Epoch [4/10], Step [370/370], D Loss: 0.9055, G Loss: 42.4359
Validation Results - L1 Loss: 0.0800, SSIM: 0.2707
Epoch [5/10], Step [1/370], D Loss: 0.8789, G Loss: 41.9285
Epoch [5/10], Step [101/370], D Loss: 0.8301, G Loss: 37.3074
Epoch [5/10], Step [201/370], D Loss: 0.8663, G Loss: 36.6406
Epoch [5/10], Step [301/370], D Loss: 1.0479, G Loss: 39.5751
Epoch [5/10], Step [370/370], D Loss: 0.8986, G Loss: 41.4412
Validation Results - L1 Loss: 0.0756, SSIM: 0.3513
```

```
Epoch [6/10], Step [1/370], D Loss: 1.0798, G Loss: 38.1668
Epoch [6/10], Step [101/370], D Loss: 1.1846, G Loss: 38.6542
Epoch [6/10], Step [201/370], D Loss: 1.0926, G Loss: 38.5350
Epoch [6/10], Step [301/370], D Loss: 0.8998, G Loss: 39.4373
Epoch [6/10], Step [370/370], D Loss: 1.1667, G Loss: 39.6570
Validation Results - L1 Loss: 0.0753, SSIM: 0.3149
Epoch [7/10], Step [1/370], D Loss: 1.0604, G Loss: 36.5707
Epoch [7/10], Step [101/370], D Loss: 1.2494, G Loss: 41.8965
Epoch [7/10], Step [201/370], D Loss: 1.2136, G Loss: 32.9834
Epoch [7/10], Step [301/370], D Loss: 1.0631, G Loss: 34.9567
Epoch [7/10], Step [370/370], D Loss: 1.0102, G Loss: 38.4644
Validation Results - L1 Loss: 0.0778, SSIM: 0.3572
Epoch [8/10], Step [1/370], D Loss: 1.4380, G Loss: 37.0367
Epoch [8/10], Step [101/370], D Loss: 1.0286, G Loss: 33.9818
Epoch [8/10], Step [201/370], D Loss: 1.1444, G Loss: 34.3082
Epoch [8/10], Step [301/370], D Loss: 1.1996, G Loss: 35.8631
Epoch [8/10], Step [370/370], D Loss: 1.3167, G Loss: 37.1101
Validation Results - L1 Loss: 0.0741, SSIM: 0.3730
Epoch [9/10], Step [1/370], D Loss: 1.0472, G Loss: 36.7639
Epoch [9/10], Step [101/370], D Loss: 1.2753, G Loss: 38.6405
Epoch [9/10], Step [201/370], D Loss: 1.2802, G Loss: 36.4594
Epoch [9/10], Step [301/370], D Loss: 1.0957, G Loss: 35.6124
Epoch [9/10], Step [370/370], D Loss: 1.2364, G Loss: 33.5947
Validation Results - L1 Loss: 0.0745, SSIM: 0.3703
Epoch [10/10], Step [1/370], D Loss: 1.0817, G Loss: 37.9398
Epoch [10/10], Step [101/370], D Loss: 1.1042, G Loss: 33.1013
Epoch [10/10], Step [201/370], D Loss: 1.0150, G Loss: 39.6250
Epoch [10/10], Step [301/370], D Loss: 1.0480, G Loss: 36.0761
Epoch [10/10], Step [370/370], D Loss: 0.9759, G Loss: 38.8151
Validation Results - L1 Loss: 0.0730, SSIM: 0.3871
```

# Experimental Results

Validation Results –  
L1 Loss: 0.0671  
SSIM: 0.4375



# Experimental Analysis


## Loss

- Discriminator Loss
  - The discriminator appeared to fluctuate a loss from roughly 0.7 – 1.2, indicating the discriminator was very effective with discrimination
- Generator Loss
  - The generator appeared to fluctuate a loss from roughly 30 - 41
  - Could be a result of poor balance between the loss weights and the discriminator and generator complexity
  - Also, could be a result of a smaller dataset, with less training epochs
- L1 loss maintained a very low number
- SSIM did not have the greatest results, which ties into the potential challenges with the generator
- Ultimately, the model generated an image with relatively similar color; however, the colors were not as vibrant and were slightly off in certain regions

## For Hinge GAN

Hinge Gan is a combination of the hinge loss function and cross-entropy loss. It performs slightly worse by the SSIM metric, but visually, colors appeared significantly more muted

- Validation Results - L1 Loss: 0.0675, SSIM: 0.4165




# Conclusion and Future Work

Was (under)stated in Anwar et al., but image colorization has a problem with performance metrics misaligning with subjective results.

Some models that we trained had similar SSIM results, but one model's outputs were significantly more saturated than the other model's and looked much better because of it.

Models tend to take the average of the choices it has seen, rather find the most likely choice (L1 Loss)



# Who did what?

- Find and preload dataset
  - Aidan
- Build pretrained model
  - Aidan
- Build decoder
  - Aidan
- Build discriminator and generator
  - Sean
- Build metric calculator and evaluator
  - Sean
- Build result visualizer
  - Sean
- Train and hyper tune model
  - Both

# Video Demo

ML-Image-Colorization.ipynb

colab.research.google.com/drive/1DqCrxdmRV8cGgiE4pYrBl9OV8Y98aCv#scrollTo=crss\_ARRx98

ML-Image-Colorization.ipynb

Commands + Code + Text Run all

Download the needed datasets

```
[1] wget http://images.cocodataset.org/zips/train2017.zip  
wget http://images.cocodataset.org/zips/val2017.zip  
unzip train2017.zip  
unzip val2017.zip  
Show hidden output
```

Import needed libraries

```
[12] import os  
import numpy as np  
import random  
  
import torch  
import torch.nn as nn  
import torchvision.models as models  
import torch.nn.functional as F  
import torch.optim as optim  
from torch.utils.data import Dataset, DataLoader, Subset  
from torchvision import transforms  
from torch.utils.tensorboard import SummaryWriter  
  
from PIL import Image  
  
from skimage import color  
from skimage.metrics import structural_similarity as ssim
```

Load the COCO datasets. Reshape to use LAB color codes: L axis in input data, A and B is output. Split train-test data here

```
[13] class CocoColorizationDataset(Dataset):  
    def __init__(self, image_dir, image_size=256):  
        self.image_paths = [  
            os.path.join(image_dir, f)  
            for f in os.listdir(image_dir)  
            if f.lower().endswith('.jpg', '.png')  
        ]  
        self.transform = transforms.Compose([  
            transforms.Resize((image_size, image_size)),
```

Variables Terminal Terminal

8:30 PM A100 High-RAM (Python 3)

# References

- Larsson, G., Maire, M., Shakhnarovich, G.: Learning representations for automatic colorization. In: European Conference on Computer Vision, pp. 577–593. Springer (2016)
- Goni, O., Arka, H. S., Halder, M., Shibly, M. M. A., & Shatabda, S. (2025, March 24). *HINGERLC-gan: Combating mode collapse with Hinge loss and RLC regularization*. arXiv.org. <https://arxiv.org/abs/2503.19074>
- Anwar, S., Tahir, M., Li, C., Mian, A., Li, C., Khan, F. S., & Muzaffar, A. W. (2024, September 2). *Image Colorization: A Survey and Dataset*. Image colorization: A survey and Dataset. <https://arxiv.org/html/2008.10774v4>
- Zanforlin, M., Munarett, D., Zanella, A., & Zorzi, M. (n.d.). *SSIM-based video admission control and resource allocation algorithms*. Second International Workshop on modeling. <https://dl.ifip.org/db/conf/wiopt/wiopt2014/ZanforlinMZZ14.pdf>